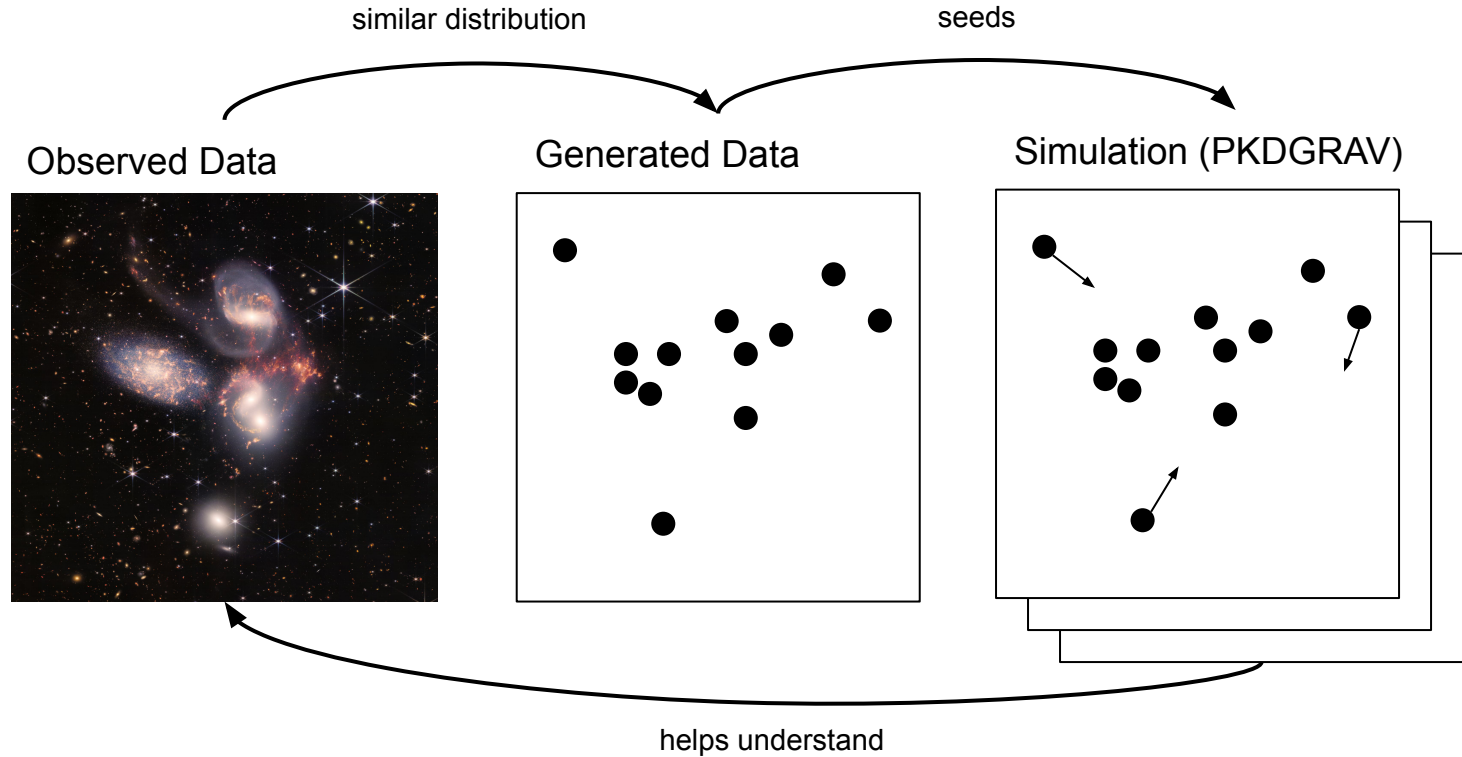



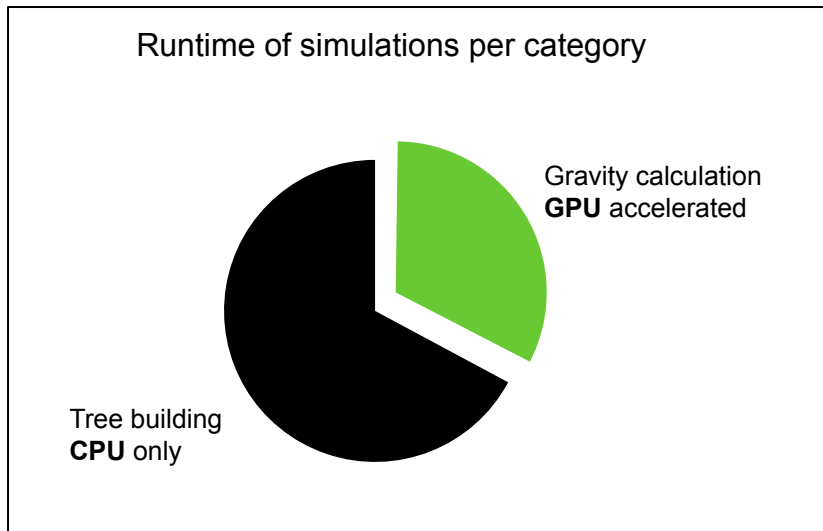
Orthogonal Recursive Bisection on the GPU for Accelerated Load Balancing in Large N-Body Simulations

Bachelor Thesis - Andrin Rehmann
UZH 2022

Astrophysical Simulations




Motivation:




As of 2020, more than $\frac{2}{3}$ of all supercomputers are CUDA enabled

Fast Multipole Method (FMM)

Naive brute force method $O(N^2)$





Fast multipole method $O(N)$




The grouping of particles is done with the **ORB** algorithm


Orthogonal Recursive Bisection **ORB**



Partition particles for fast access




Bisection method






We aim for $12 / 2 = 6$
particles smaller or
left of the cut

$O(32)$ runtime for 32 bit precision

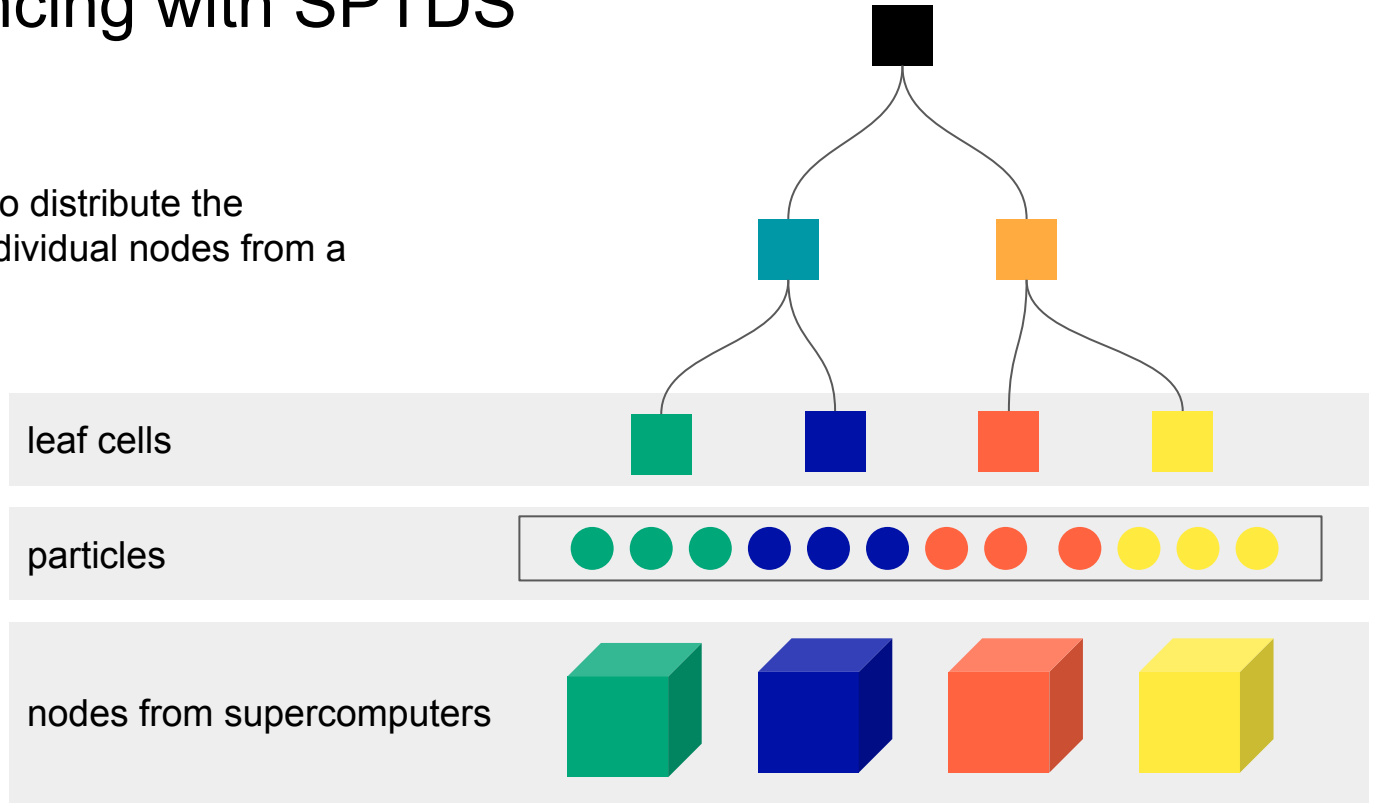


Space Partitioning Tree Data Structure (SPTDS)





Load Balancing with SPTDS

Tree can be used to distribute the particles among individual nodes from a supercomputer




6.2x speedup of GPU over CPU version expected





1. Use GPU as **much** as possible
2. Use data link between CPU and GPU as **little** as possible

CUDA




- thousands of blocks can exist
- up to 1024 threads per block
- only threads from same block can access a shared memory register
- all threads can access the global Memory

Synchronization




Warps



branch divergence: thread executes different code from other threads in warp

- A warp consists of 32 threads
- Synchronization is possible and needed to avoid data race
- Warp level primitives can be used to implement most algorithms on thread level

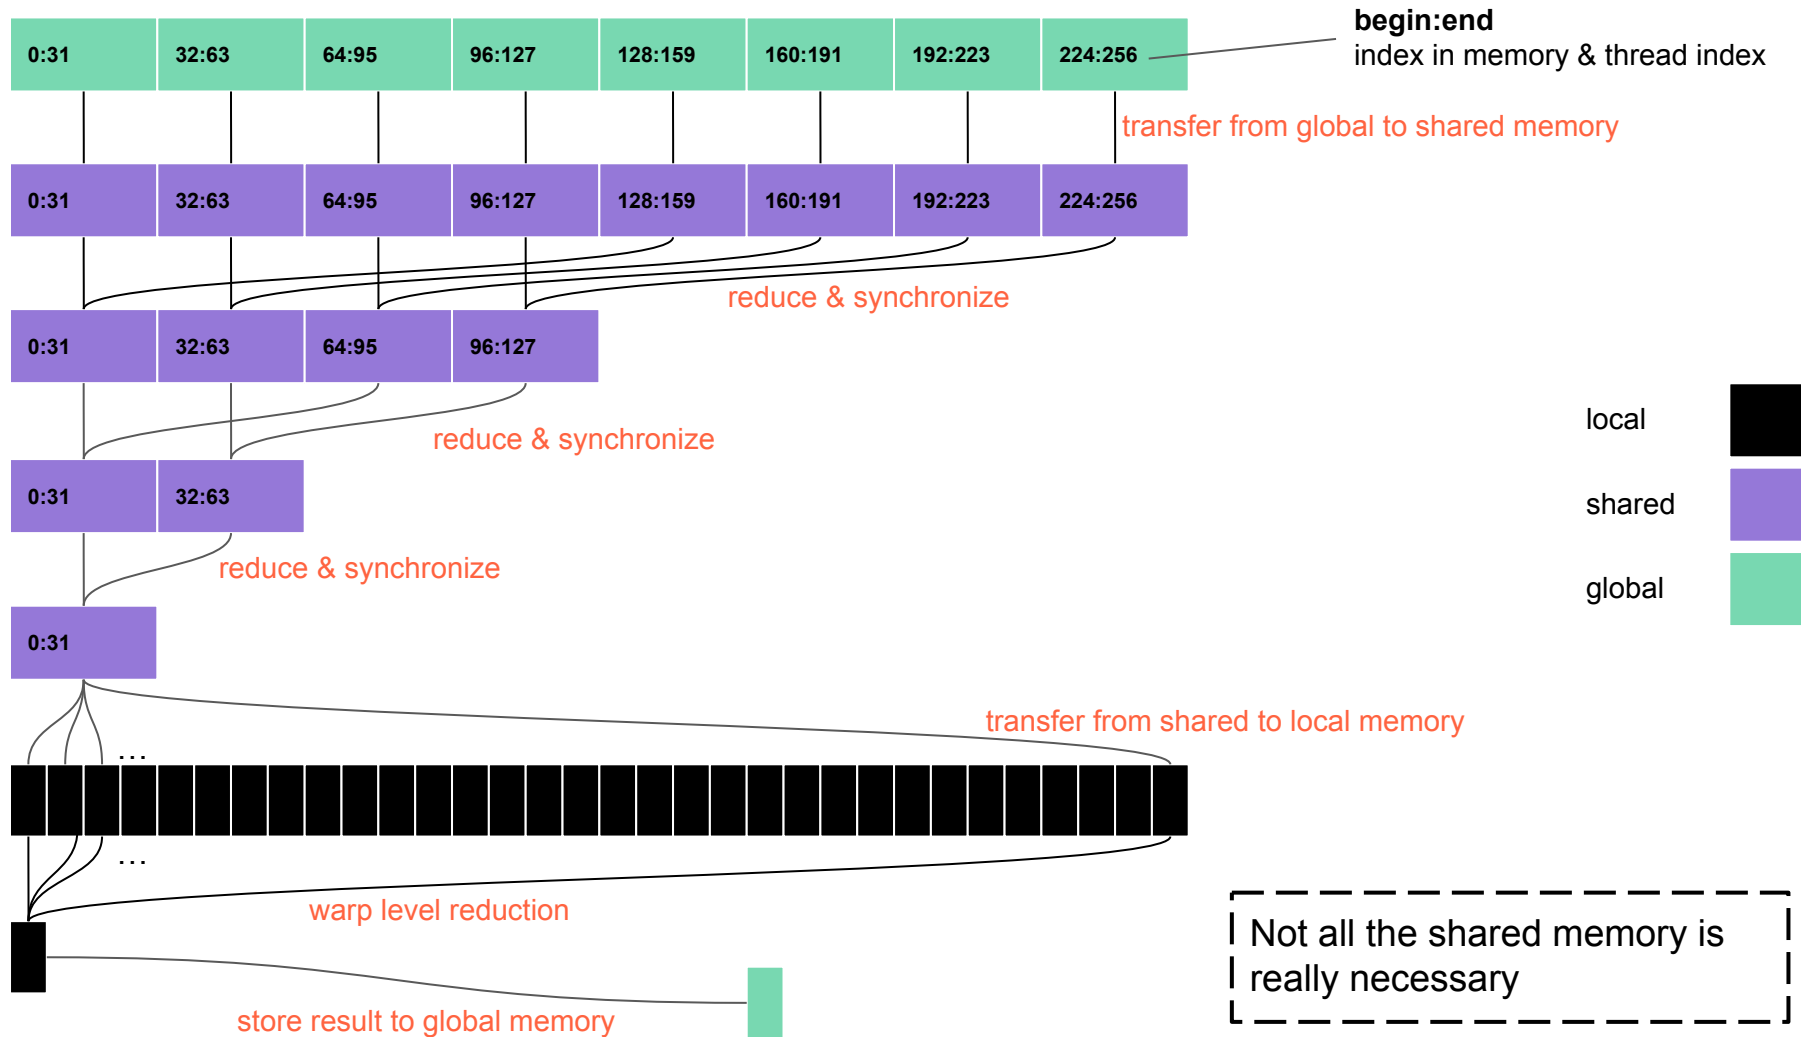
Bisection method



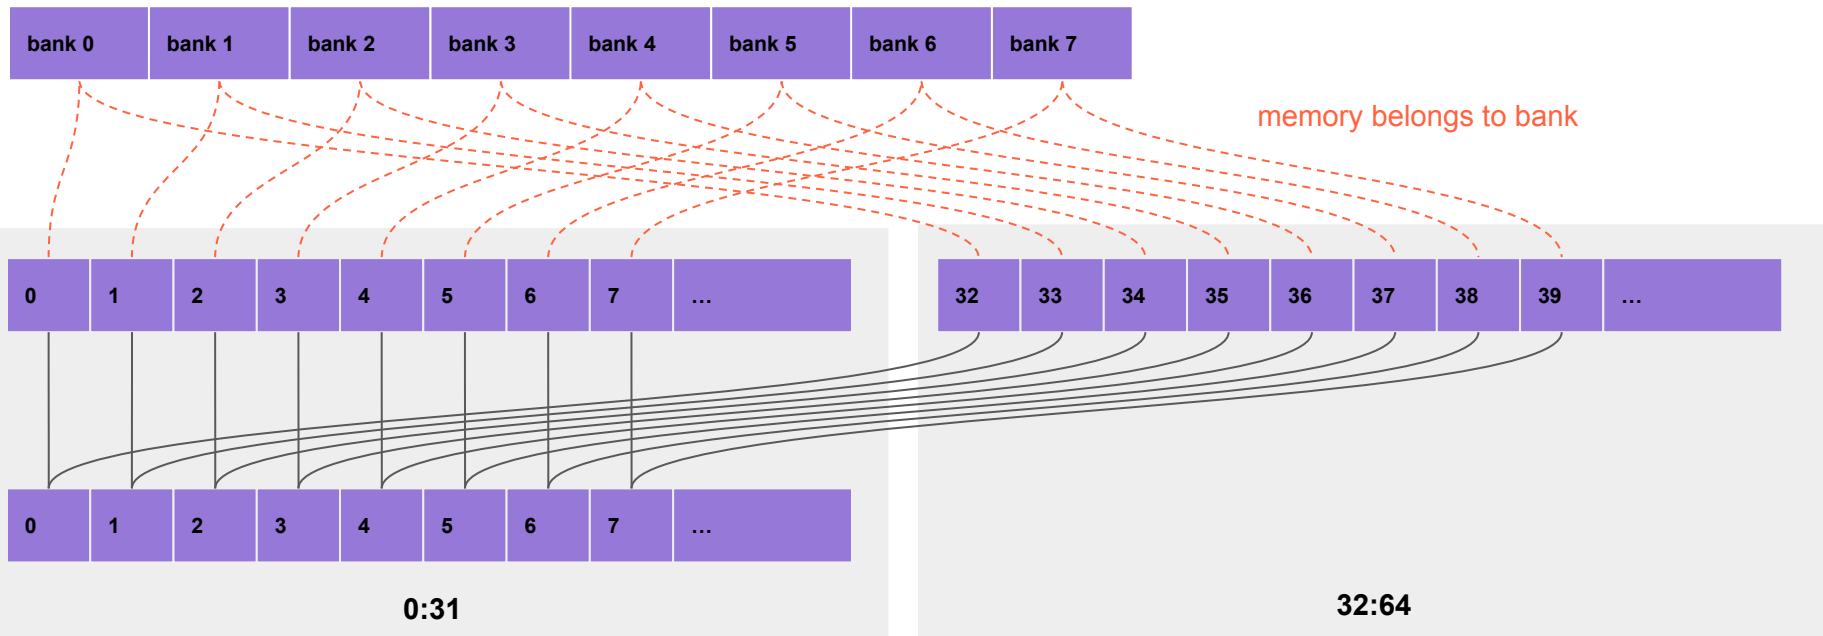
$$0.9 < 0.5 + 0.1 < 0.5 + 0.3 < 0.5 + 0.6 < 0.5 + 0.7 < 0.5 + \dots$$

$$(0.9 < 0.5 + 0.1 < 0.5) + (0.3 < 0.5 + 0.6 < 0.5) + (0.7 < 0.5 + \dots)$$

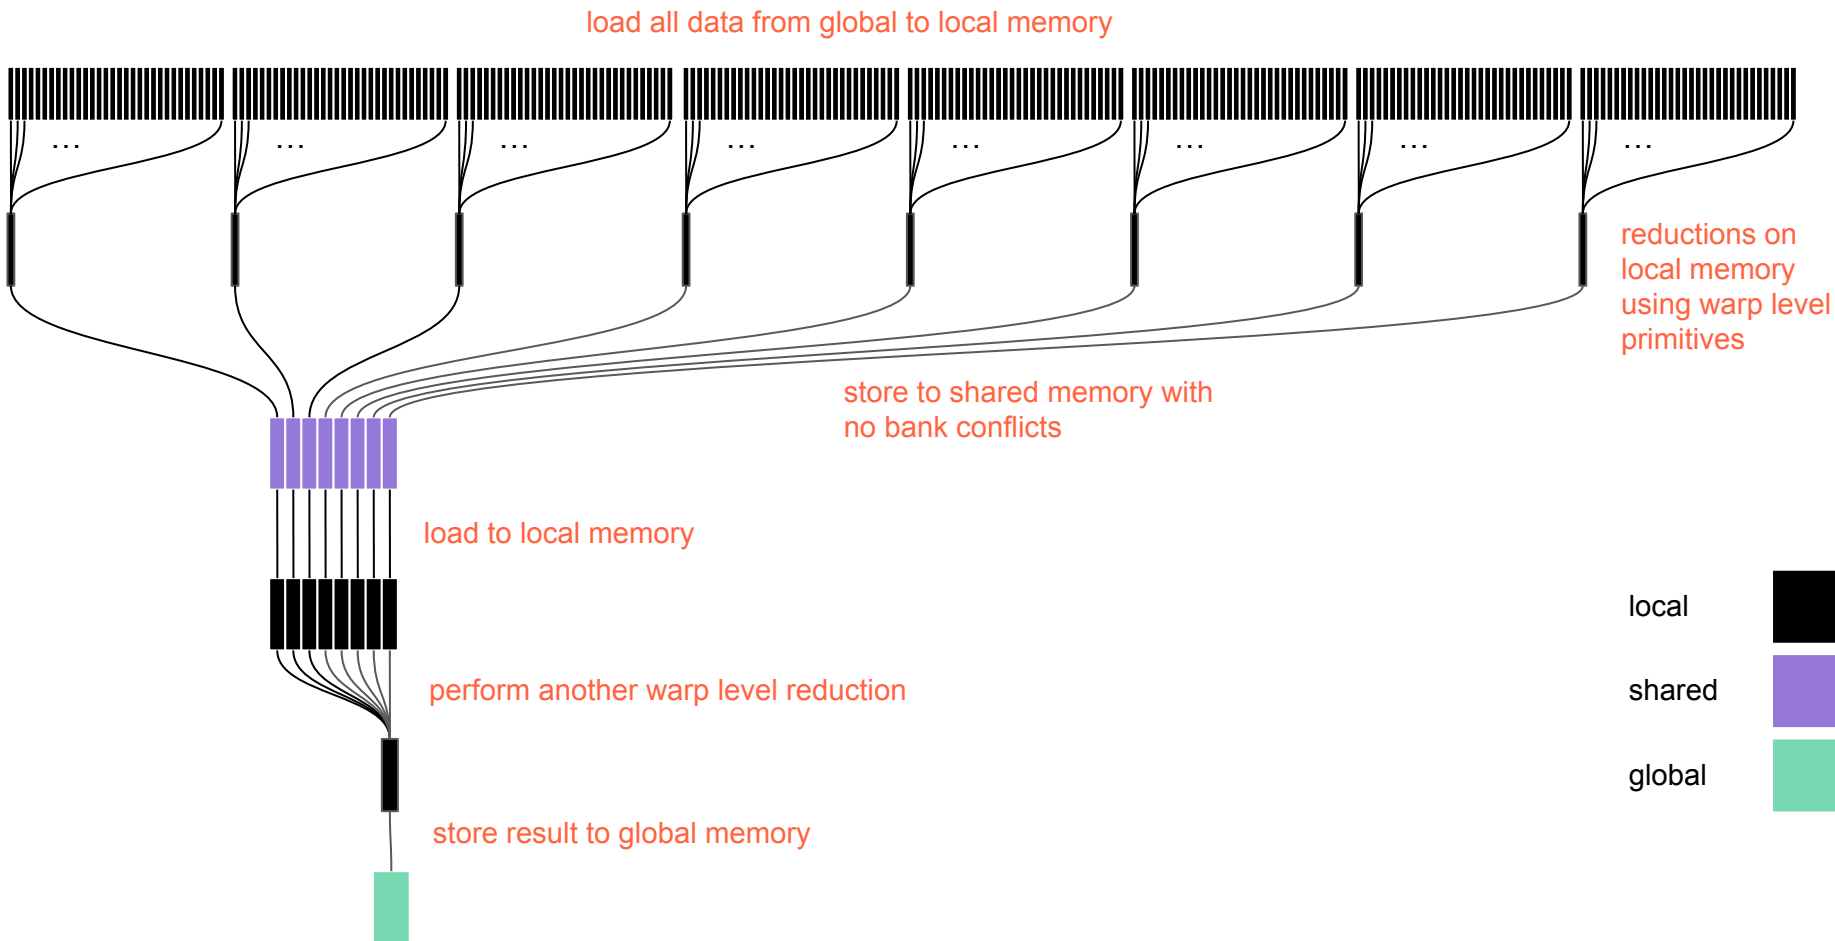
We can easily split the Count Left part of the bisection into multiple sub problems: ideal for CUDA.



Bank conflict free access pattern



each thread of the same warp only accesses a single memory bank



particles stored in global memory



pointers to particles passed
as parameter to kernel

cell info passed as
parameter per kernel



Kernel #1



Kernel #2



Kernel #3



Kernel #4

each block performs block wise reduction

number of blocks per kernel might vary


particles stored in global memory



Single pointer passed to kernel

Cell information is
copied to shared
memory in block wise
entries

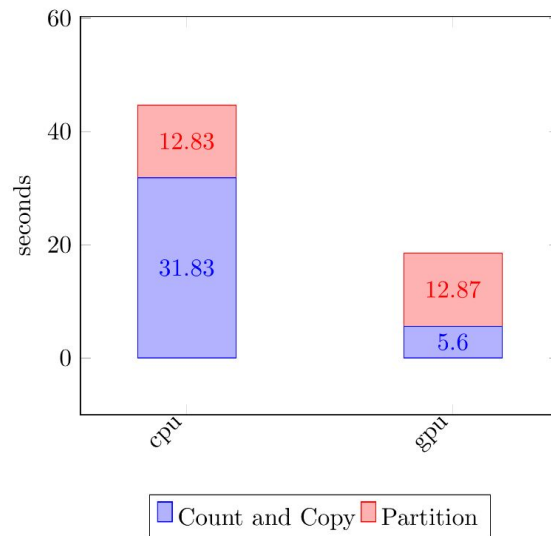
Each block reads
assigned info from
global memory



Kernel #1

Conclusion

- Successfully improved runtime of ORB 🎉
 - **5.6x** speedup. Theoretical limit was **6.2x**
 - Runtime also close to theory in absolute terms
- Better understanding of hardware
- Learned CUDA
- Improved C++ skills



Outlook

- Integrate to PKDGRAV
- Improve CPU partition
- Finish implementing GPU partition
 - ◆ not clear weather could give improvements
 - ◆ higher memory usage
- Data compression

